

Kitap : Modern OS Tonenbaum and Bas 4<sup>th</sup> edition.

Yüzdeker (Approximately) : 01030 midterm, 01040 final, 01030 HW 01.03.2021

## OPERATING SYSTEMS

121 fsh

OS is a piece of software.

OS works on HW

HW (CPU+memory+Disk)

\* That abstracts the computer HW through interfaces.

HW abstraction

- Hides the many details underlying HW
- Prevent users with a resource abstraction that is easy to use
- Extends or virtualizes the underlying machine

\* Manage the resources

Resource management

- Processors, memory, timers, disc
- Allows multiple users and programs to share the resources and coordinates

OS is managing who is running now (timer, memory, syscall...)

OS → kernel mode

\* Performs services

- Abstract HW
- Provides protection (gcc, ash memory environment etc.)
- Manages resources

MSKII → m interrupt (Kernel'da interrupt olmasa crash yapabiliriz.)

As RM, sharing resources in 2 different ways

1- In time (Time sharing ex. printer, CPU, Network Adapter)

2- In space (Disc, Memory)

- α allows multiple programs run at the same time
- α Manage and protect memory

Without OS, we can't run so much HW

## OS vs. Kernel

### • Wide view

Windows, Linux are OS

Includes system programs, libraries...

### • Narrow Definition

OS often equated with kernel

The Linux kernel; the Windows executive - the special piece of sw that runs with special privileges and actually controls the machine.

2.3.2021

OS not always running.

ls, ps etc. is part of the standard bash shell

So they are not part of the OS.

\* For Linux, drivers are in kernel mode.  
memory management      "      "  
signals                      "      "

## History of OS

### 1 - 1<sup>st</sup> Generation (1945-55)

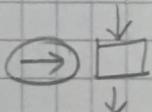
Vacuum Tubes.

Her 20 dokikada bir "pathogen" tipler olduğu için verimiz.

### 2 - 2<sup>nd</sup> Generation (1955-65)

Transistor and Batch System

Transistor controls the flow of current



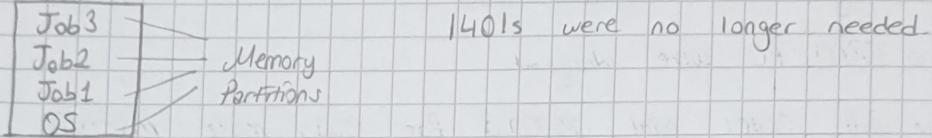
Punchcard : 80 columns 12 rows cards

cards → card reader → computer (1794) → printer (1401)

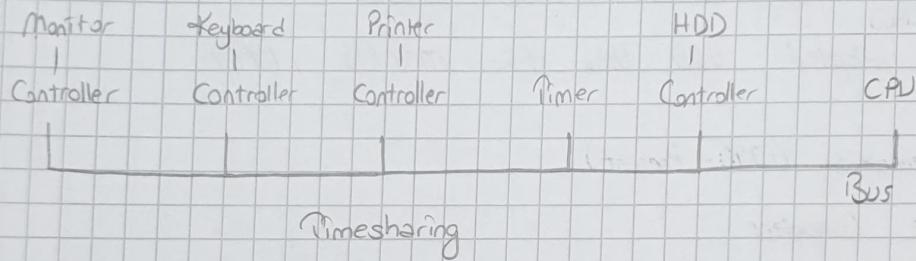
Computer was very expensive

3 - 3<sup>rd</sup> Generation (1965-1980)

IC's and Multiprogramming



I/Os were no longer needed.



Still cost of million dollars

#### - Multiprogramming

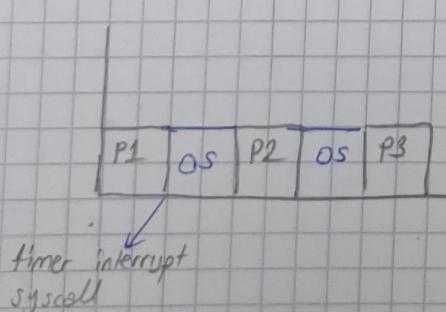
Run several programs at once

Protect each one's memory against the others

When one program needs to do I/O, let another program uses the CPU

As necessary pre-empt, a CPU bound program to let another run

Use OS "daemons" for queuing → *değin sırıza çoklu plan sistem*



## 4- 4<sup>th</sup> Generation (1980- Present)

### Personal Computers

- \* Large Scale Integration Unit
- \* IBM PC were designed
- \* Bill Gates paid 75000\$ for DOS, made MS-DOS
- \* Engelbart invented GUI, mouse, window
- \* Apple adopted BSD Unix for macOS , 1999

## 5- 5<sup>th</sup> Generation (1990- Present)

### Mobile Computers

- \* Nokia released N900 (mid 1990)
- \* Ericsson coined term smartphone (1997)
- \* RIM's Blackberry OS (2002)
- \* Apple iOS (2007)
- \* Android Linux based OS by Google (2008)

Security and battery management

4th Generation (1971 - Present)

Large Scale Integration Circuit  
IBM designed the IBM PC

Bill Gates paid \$5000 for DOS, made MS-DOS

Eugelbart invented GUI, mouse, windows

Apple adopted BSD Unix for macOS, 1999

5th Generation (1990- Present)

Nokia released N9000 (mid 1990s)

Ericsson coined term smartphone (in 1997)  
RIM's BlackBerry OS (in 2002)  
Apple iOS 2007

Android Linux based OS by Google 2008

Security and battery management

8.3.21

MMU → Memory Management Unit

Why are we Responsible for making virtual memory possible. Without VM, OS doesn't work. makes our processor think that ~~it has always~~ it has always.

HD controller has register to read and write

### Processor

Brain of the computer  
Fetch, decode, execute  
CPUs have registers  
PC

Stack Pointer

Program Status Word (PSW) : condition bits, kernel mode, CPU priority  
A user program must take a syscall, which traps into the kernel and make OS

pipelining

flag checking enables

global doesn't use stack

gcc -f test.c  
less test.c

### Pipelining

F → D → E

F → D → Holding Buffer → E

8.03.21

## \* Multithreading / Multicores

Moore's Law  $\rightarrow$  x2 processors every 18 months

Single processor can process more than one thread efficiently (multithreaded)

Multicores: Real multiple CPUs on the same chip

GPUs are extreme examples of multicores (thousands of cores)  
simple than CPU

## \* Memory

Should be fast, large, cheap

Not all possible with the current technology

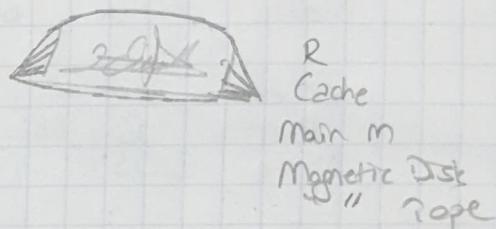
## \* Registers and Cache

(register)  $32 \times 32$  bits for  $32$  bit system  
 $64 \times 64$  " in  $64$  " "

No delay accessing them

L1 cache 16 kB

L2 cache a few megabytes



## \* Disks

Larger, cheaper, slower

Slow because they mechanical

spinning

Load OS and apps from disk

Manage disk space allocation

Store and retrieve files by name

Provide some protection for files

## \* I/O Devices

Controller and the device  
Very complicated, controller handles this  
Driver: Middle layer SW for device controller by manipulation

Three types of I/O

1- Direct memory access  $\rightarrow$  CPU, Disk writer (I/O is over)

2- Interrupts  $\rightarrow$  CPU asks the driver to get the data blocks on int. When data is ready, it is transferred to the device

3- Direct Memory Access (DMA)  $\rightarrow$  CPU sets up the DMA chip

telling it how many bytes to transfer, the device and memory addresses involved, and the direction, and lets it go

1-Interrupt  
2-Dispatch  
3-Given

## Interruptions vs traps

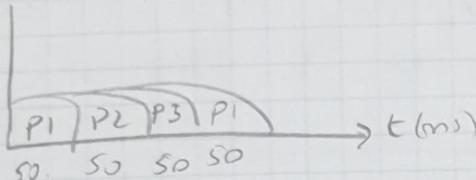
(HW) ↓  
1-way

(SW)

- 1- push byte buffer
- 2- " " " " " " " "
- 3- " " " " " " " "
- 4- syscall
- 5- " " " " " " " "
- 6- Put code for read in reg.
- 7- Dispatch syscall handler
- 8- Syscall hand
- 9- Return to caller
- 10- Increment SP

9, 3, 21

Thread context switch costs much



lets 1000 1000 1000 1000 → bad interactivity, less switch

OS 20

Mainframe, Server, Mult. Processor, PC, Embedded, Handheld, server node

Server OS

Provide server some of them large. It can provide print service, file service  
Typical are Solaris, FreeBSD, Linux, Win 201x  
Some multiple users

~~Mainframe  
server  
large  
distr~~

Smart card OS

JVM  
Credit card

Smallest OS  
ROM holds an interpreter

Multiprocessor OS

More than one processor

Job of OS is more difficult  
the hard part

PC

Linux, Win, Apple  
Interactivity important

Handheld

Android, Apple iOS smartphone and tablets

Embedded

run on computers that control devices  
TV's, microwave

No need for protection

Don't accept user-installed software

Server Node OS

Is a real comp with CPU, RAM  
ROM

Small, real OS, event driven

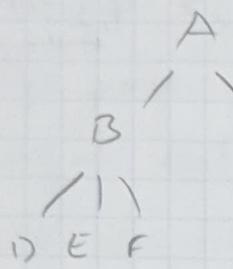
Real-time OS

Must occur at a certain moment (or range), we have a hard real time sys  
There is not a requirement to wait  
50 ms

Multiprocessing  $\rightarrow$  Doesn't need multi processing

## Processes

①



pid = fork()  
return 0 if it is child.  
→ child process. pid = waitpid(pid, &status, options)  
s = execve(name, argv, envirnp)  
PS -0xf  
exit()

②

## OS Concepts

1- Processes

Address Spaces

Files

I/O

Protection

The Shell

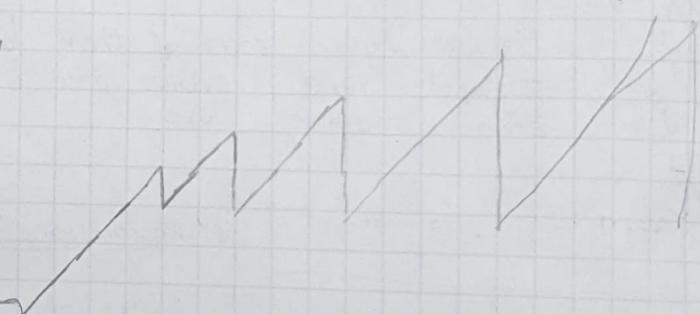
Integrity recapitulates phylogeny

Large memory

Protection hw

-Disks

-Virtual memory



## 1- Processes

It is basically a program in execution

Process table, array of structures, one for each process currently in existence

- Registers, open files, other data

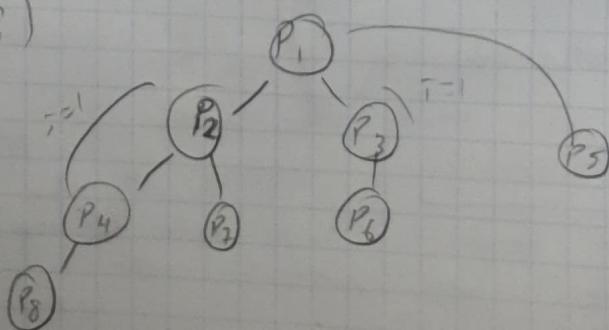
If parent is terminated, OS kills child process.

for (i=0; i<5; ++i)

fork()

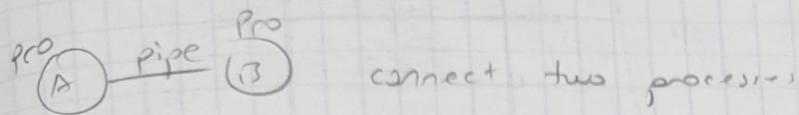
$2^{5-1}$

$1+2+4+8+16$



## Files

most PC OS have the concept of a directory as a way of grouping files  
 - Path name, root directory, working dire., file descriptor



ls -alRt | wc -l

## Integry --

Development of an embryo (ontogeny) repeats (rec...) the evolution of the species (ph.)

Each new species of computer

PCSI X has about 100 procedure calls

" proc calls onto system calls is not one-to-one  
 OS calls are too expensive (user to kernel, data, etc)

Proc. and syscalls are different

exit has no path. But it runs flaw? if (strcmp(command, "exit") == 0) exit(0)

16.03.2021

## syscalls for File Management

fd = open(file, how, ...)

s = close(fd)

n = read(fd, buffer, nbytes)

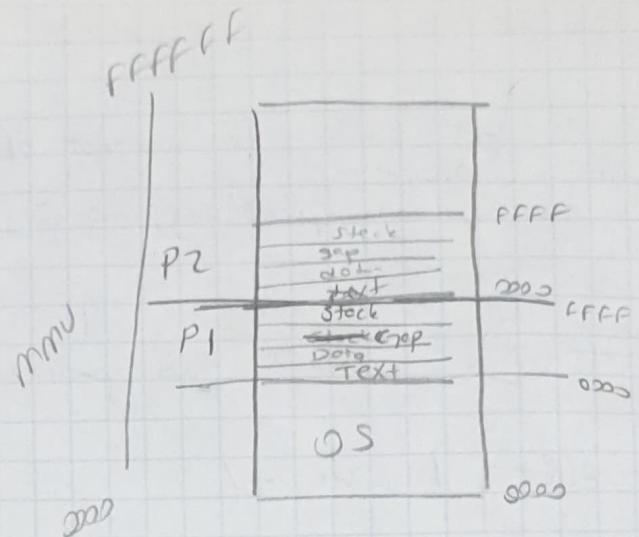
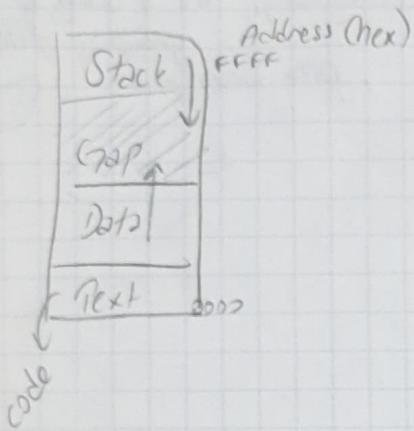
n = write

sends signals

link(name, name<sup>2</sup>)  
 unlink(name)  
 mount(special, name, flag)  
 umount(special)

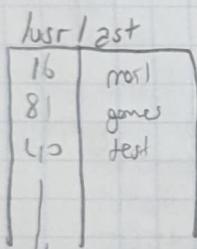
s = chdir(dirname)  
 = chroot(dirname, mode)  
 s = kill(pid, signal)  
 seconds = time(&seconds)  
 Miscellaneous Syscalls

## Memory Layout

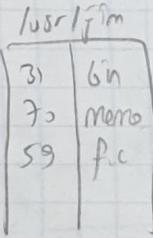


## Linking

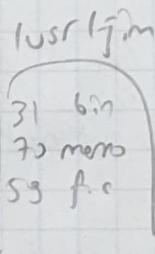
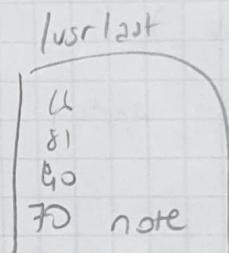
object  
i-node



(a)

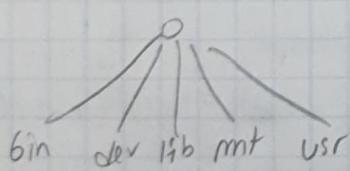


S checksum  
word 13  
ptr 16 25 50  
13

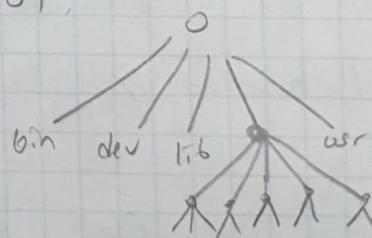


link("usr/lib/memo", "usr/last(note")

## Mounting

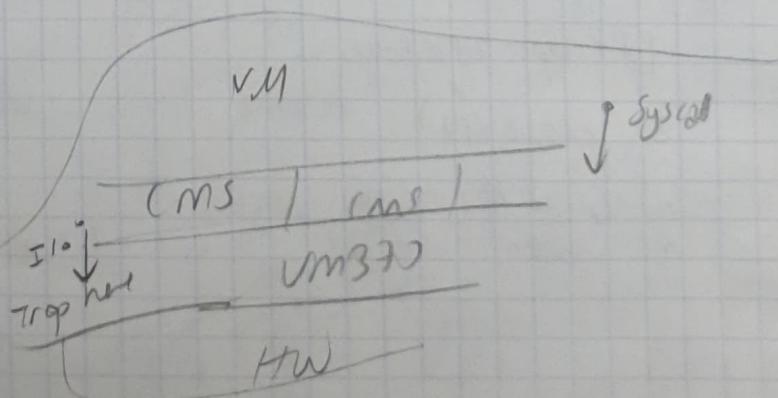


mount("/dev/sdb0", "/mnt", 0);



## Miscellaneous Syscall

chmod -rwx linked.txt

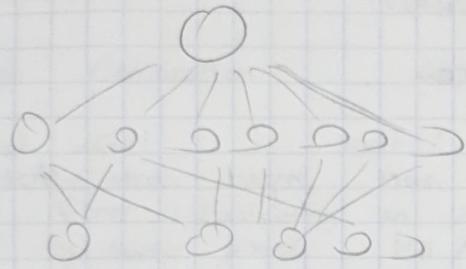


The JVM  
Disadvantage → Slow  
Ex kernels  
OS Partitioning

If real in VM it wants reach HW, VM can handle op.s. but process we have supervisor  
hypervisor

## OS Structure

### • Monolithic Systems



Main Procedures

Service procedures (file I/O etc.)

Utility Process

### Layered Systems

ancestor of UNIX

Instead of layers, MULTICS was described as having a series of concentric rings.

- 5 Operator
- 4 User proc
- 3 I/O
- 2 Op-process com.
- 1 Memory and drum man.
- 0 Processor allocation



### Microkernels

1000 - 1.3UG

Monolithic OS of 5 million line codes is contain between 10k-50k kernel bips.  
Keep kernel small!

MINIX 3 microkernel is 12.000 lines C and some 1400 lines assembler

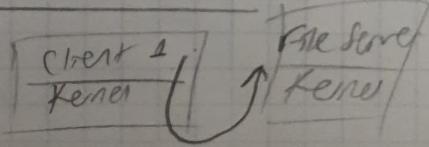
- Policy user mode (All algorithms called <sup>simplified</sup> policy)  
 - Mechanism kernel (Reaching high) (Not smart)
- Microkernel handles Interrupts, processes, scheduling, IPC

Other of them is in user mode

Overclocking : Increasing the clock rate of CPU 1 GHz  $\rightarrow$  1.2 GHz

Drivers in user mode is smart. Not crash whole OS

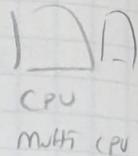
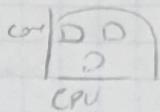
### Client-Server Model



## CHAPTER 2:

### PROCESSES AND THREADS

Multicore

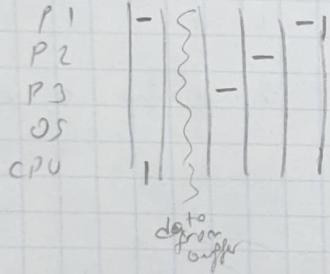


- Most central concept in any OS is process:  
An abstraction of a running program

- Pros are one of the oldest and most import abstr. that OS provide

### Single Program vs. Multiprogramming

Each pro runs to completion, but intermixed with other pros..



16-03-2021

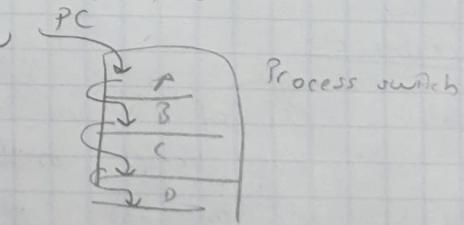
### Process Model

A process is just an instance of an executing program

Conceptually each process has its own virtual CPU

We will assume there is only one CPU

Increasingly however that assumption is not true, since new chips are often multicore



Multiprogramming

### Process Creation

- 1- System initialization; Processes that stay in bg to handle some activity  
mail etc and so on are called **daemons** → long running process that handles requests for service
- 2- Execution of a process creation syscall by a running process (fork)
- 3- A user request to create a new process
- 4- Initialization of a batch job

Create new pro fork  
change its memory img execute  
Windows 100 commands

### Process Termination

- 1 - Normal Exit      exit / ExitProcess
- 2 - Error Exit
- 3 - Fatal Error : Executing an illegal instruction, referencing nonexistent memory, divide by 0
- 4 - Killed by another process Kill / Terminate Process

### Process Hierarchies

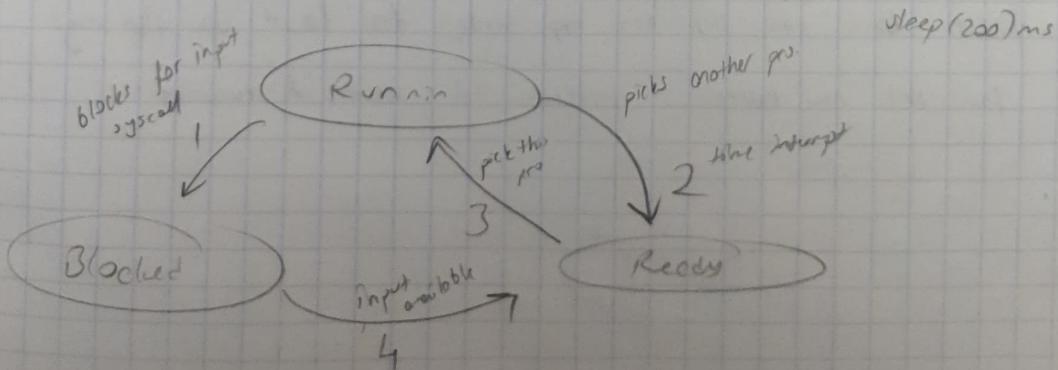
- In some systems, parent and child continue to be associated in certain ways
- In UNIX process and all of its children together form a process group. A special process, called init, is present in the boot image
- Windows has no concept of a process hierarchy

### Process States

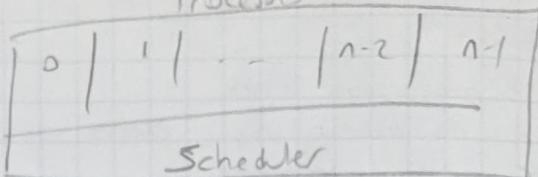
cat chapter1 ch2 ch3 | grep tree

Depending on the relative speeds of 2 process, it may happen that that grep is ready to run, but there is no input waiting for it. It must then block until some is available.

- (1)ode) 1 - Running (using CPU at that instant)
- (several) 2 - Ready (runnable, temporarily stopped to let another process run)
- 3 - Blocked (unable to run until some external event happens)



## Process



Lowest level of OS is scheduler

All the interrupt handling and details

of actually starting and stopping process are hidden away in what is here called sched.

## Impl. of Process

A table (an array of structures) called process table, with one entry per process

Process Manager	Mem. Man	File M
Registers	Pointer to -- reg.info	
PC	stack	Root directory
SP	data	Working "
PID	text	File descriptor
CPU time used		User ID
Alarm		Group ID

User process is running when a disk interrupt

- 1- HW stacks PC, prog. status word (not all registers, why?)
- 2- HW loads new PC from interrupt vector

HW saves only  
minimal register  
PC or status address own  
over one operation

- 3- Assembly - lang. procedure saves registers. All interrupts saving the registers
- often in the ptable entry for current process

- 4- " " sets up new stack saving the registers and setting the stack pointer cannot even be expressed in C!

5- C interrupt service runs (typically needs ad buffers input)

6- Scheduler decides which process is to run next

7- C procedure returns to asm code

8- Control is passed back to the asm loop code to load up the regs

and memory map for the new current process and start it running

Address of main.main() buffer  
handling of interrupt register

Interrupt  
flow

## Modeling Multiprogramming

P<sub>1</sub>      300ms wait I/O      → Efficiency 70%

P<sub>2</sub>      300ms wait      →

CPU utilization =  $1 - p^n$       in process are waiting for I/O is  $p^n$

P <sub>1</sub>	RU	B <sub>1</sub>	REIB	
P <sub>2</sub>	BRE	B	RU	

$0,3$   
 $\times$   
 $0,3$   
 $= 0,09$

$$1 - 0,09 = 91\% \text{ Efficiency}$$

For 3 pro.       $1 - (0,3)^3 = 0,1098$

Ex

OS takes 2GB  
Each user prog 2GB

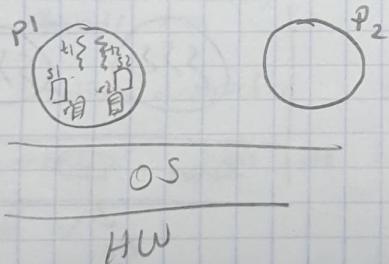
in 8GB with 0.680 wait  
we can run 3 programs at once

$$\text{CPU ut.} = 1 - (0,18)^3 = 49\%$$

## THREADS

Like processes: programs in parallel

Share the same address and data

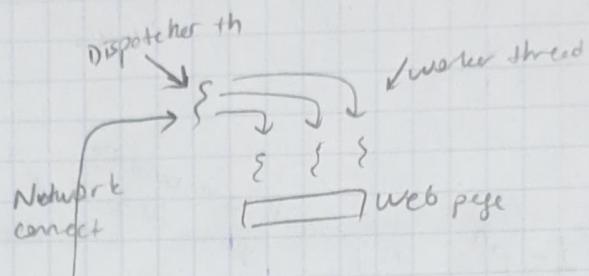


- Lighter and faster to use
- There is no protection b/w threads
- Threads co-operate, they do not compete

< Main reason is in many app., multiple activities are going on at once  
 Some of them may block from time to time

for ex. → MS Word has 3 threads

" " → Web Server



### Dispatcher Thread

while TRUE

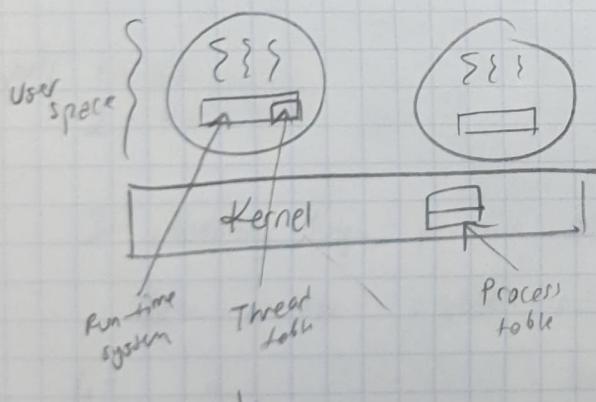
get next request (lbuf)  
 handoff work (lbuf)

### Worker Thread

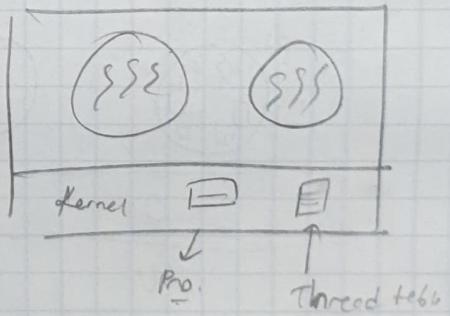
while TRUE

wait for work (lbuf)  
 look for page in cache (lbuf, lpage)  
 if page not in cache (lpage)  
**LOAD HIT** read page from disk (lbuf, lpage)  
 return page (lpage)

### Imp. Threads in User space



1<sup>st</sup>  
 OS doesn't know thread



2<sup>nd</sup>

<sup>main</sup>  
 Adv. 1 Speed no context switch  
 2 Handle everything in user (flexible)

Disadv: If one of threads is blocked,  
 all threads will be blocked

## Model

Threads	Parallel ✓	Blocking system
Single Thread Processes	No parallel	Block
FSM Machine	Parallel	NonBlock

7

Lightweight processes

If one thread open a file, others can access them

Each thread has own PC, Registers, Stack, State

23.02.2021

Kernel threads → Unix, Windows etc

In modern OS → Parallelism is important (Run threads ~~at~~ in same time)

## User Level Threads

- Thread creation, destruction, scheduling done at user level
- Thread scheduler finds the highest priority thread

No syscall ✓

No context switch ✓

No copying data safely across protection boundary ✓

No need to flush the memory cache ✓

Allow customized scheduler

## Disadvantages

Blocking syscall

Page faults

Inability to benefit from multiprocessor CPUs

Kernel can't automatically grow per-thread stacks

## Imp. Threads in kernel space

Kernel manages thr

Similar thr table but in kernel

No problem with blocking syscall

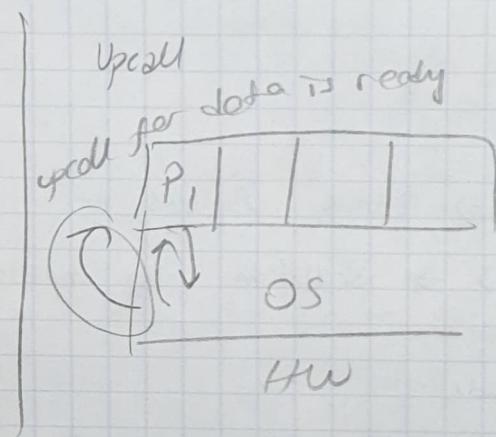
Handles mult CPU easily

Thread switching is expensive

Still cheaper than process switches

Preemptive Scheduling  $\rightarrow$  Pro. uses too much CPU, I am going to take the CPU away

## Hybrid



## Scheduler Activations

Thread creation is done at user level

When file is ready, kernel does an upcall

Careful about non-blocking

## Making Single Thread Code Multithreaded

Local variables are not cause any trouble

One solution is to prohibit global variables altogether

Another  $\rightarrow$  It is possible to allocate a chunk of memory for the globals and pass it to each procedure in thread or parameter  
create-global, use-global

Static shouldn't use

Must use reentrant routines

can be invoked more than once simultaneously

can only use local or dyn.a

must not use static

many c libraries routines do use static buffer

## Interprocess Communication (IPC)

first issue: One process can pass info to another

2<sup>nd</sup> issue: 2 process airline reservation system to grab last seat

3<sup>rd</sup> issue: Process A produced data, B reads it. B has to wait until A has produced

## Race Conditions

2 or more processes are reading or writing same shared data  
and the final result depends on who run precisely

## Critical Regions

Mutual exclusion: One process shared variable lockstep, developer will be excluded

Part of the prog. where the shared memory is accessed is critical region

Requirements to avoid race conditions:

- 1- No 2 processes may be simultaneously inside their critical reg.
- 2- No assumptions may be made about speeds or the number of CPU
- 3- No process running outside its critical region may block other processes
- 4- " " should have to wait forever to enter its critical

## Mutual Exclusion With Busy Waiting

Continuously testing a variable until some value appears

- Proposals for achieving mutual exclusion

• Disabling Interrupts

• Lock variables

• Strict alternation

• Peterson's solution

• TSL Instructions

### Disabling Interrupts

available only to the kernel

Better not be available to user process!

Doesn't work well on multiprocessor

Special case solution for kernel only

### Lock Variables

global variable lock

while (lock != 0)

Start critical

lock = 1

End

lock = 0

There is a  
window  
and check if P1  
and set it to 1

P2

source

### Autost

```
white (TRUE) {  
    while (turn != 0); // loop  
    critical_region();  
    turn = 1;  
    noncritical();  
}
```

A lock is called a spin lock

a is blocked enter region

memory  
busy

```
white (TRUE)  
white (turn != 1);  
critical();  
turn = 0;  
noncrit();
```

strict alternation

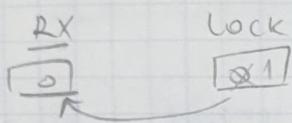
1000000 saying whether  
a random thread decriticals

## - Peterson's Solution

```

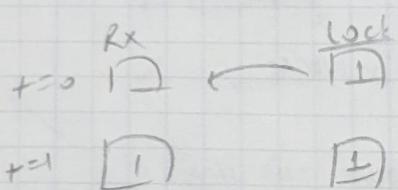
while(1)      P0
    enter region(0)
    critical region
    leave region(0)
    non critical
    →
    P1 same
  
```

## - TSL Instruction (Test and Set Lock)



locktaki değer RX'e alı, sonra lock'u değiştir.  
It's guaranteed to be indivisible

or



## - Spin Locks & Priority Inversion

I have 2 processes. Higher Lock

L never runs if H is runnable

While H sleeps, L grabs the spin lock

H wakes up and tries to get the lock

H spins, waiting for L to free it

But L can't get the CPU, so it doesn't progress

One example of deadlock

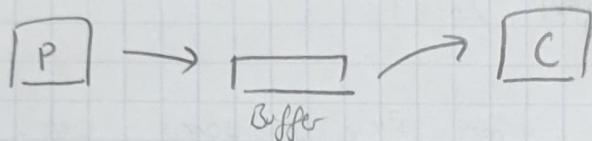
## - Blocking Instead of Busy Waiting

Prevent wasting time busy waiting

Simplest way sleep + wakeup

Syscall sleep causes to block (suspend until another process wake it up)

## Classical Problem: Producer - Consumer



## Race Condition

### Producer

```
while(1){  
    item = produce()  
    if(count == N) sleep() //empty  
    insert item()  
    if(count++ == 0)  
        wake(consumer)}
```

### Consumer

```
begin →  
if(count == 0) sleep()  
item = remove  
full ← if(count-- == N)  
wake(producer)  
consume item
```

## Semaphore

down - up

↓      ↓

sleep

wakeup

↓

decrement  
if st > 0

↓ increment

Binary semaphore = mutex 0 or 1

## Implementation

Typically done by kernel

Mask interrupts while manipulating semaphore

On multiprocessor, use TSL as well

Both of these are for only a few ms - not a serious problem

semaphore mutex = 1  
semop empty = 100  
semop full = 0

while(1)

item-produce  
down(&empty)  
down(&mutex)  
insert  
up(mutex)  
up(full)

down full  
down mutex

up mutex  
up empty

---

## 2 Different Uses of Semaphores

1- Counting semaphores (full, empty)  $\rightarrow$  Sleep / wait

2- Mutual exclusion semaphores (mutex)  $\rightarrow$  Make sure just 1 process in critical region

Mutexes are special case of semaphore. Useful when no counting

If TSL is available, easy to implement at user level

Try to grab lock, if unsuccessful let another thread run.

## Different Processes?

If they have disjoint address spaces, how can they share semaphores or a common buffer?

1- Shared vars in kernel

2- Many OSes offer a way for processes to share some portion of their address spaces with other processes.

## Futex (Fast User Space Mutex)

Busy wait or lock faster?

Block only if necessary and switch to kernel space

## Condition Variable

Second synch mechanism (Allows switch between 2 cond. latches)

Cond. var. allow threads to block due to some condition not being met

cond\_wait(&condp, &the\_mutex) → needs both of them

## Risks of Semaphores and Mutexes

Using correctly is difficult

If you get wrong, deadlock!

Testing is difficult, because it's all timing-dependent.

We need higher level construct.

## - Monitors

A pl. construct

Java has, C/C++ not

Like a class, but only 1 thread can be executing in it at a time

Like language implementation of mutexes

They are good for mutex and you cannot go wrong with mutex

Ex:

monitor example

integer i  
condition c

procedure producer()

;

end

procedure consumer()

end

end monitor

Automatic Mutex  
Automatic Handler

## Monitors and Application Blocking

We still need a way to block

2 new operations: wait and signal

### Java

```
static our-monitor mon = new our-monitor()
```

## Disadvantages of Monitors

Only available in a very few languages

They don't work well without some shared memory

Vimoliye kodor user or kernelda shared memory varsa dige konustuk

- spin locks
- semaphores
- mutexes
- cond variable
- monitors

What if memory space is not shared? Processors are running on 2 different machines

## Message Passing

```
send(dest, &msg)
```

Sender blocks if receiver's buffer is full

```
recv(src, &msg)
```

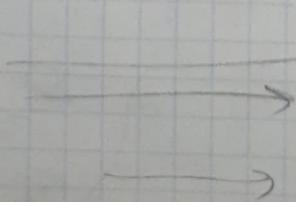
Receiver blocks if there is no data avai.

-Consumer Producer Gangi:

```
producer(void) {  
    int item  
    message m
```

while TRUE

```
    item = produce  
    receive (consumer, &m)  
    build-message (&m, item)  
    send (consumer, &m)
```



```
consumer(void)
```

```
for (i=0 i<N ) send (producer, &m)
```

```
receive (producer, &m)
```

```
item = ...
```

```
send (...)
```

```
consume-item (item)
```